

ASSIGNMENT : Bagging & Boosting KNN & Stacking Assignment

QUESTION 1 : What is the fundamental idea behind ensemble techniques? How does bagging differ from boosting in terms of approach and objective?

ANSWER 1 : The **fundamental idea behind ensemble techniques** is to **combine multiple individual (weak) models** to create a **stronger, more accurate, and more robust predictive model**.

Instead of relying on a single model, ensemble methods aggregate the predictions of many models to reduce **variance, bias, or overfitting**, depending on the specific technique used.

1. Bagging (Bootstrap Aggregating)

Approach:

- Bagging builds **multiple independent models** (usually of the same type, like Decision Trees) on **different random subsets** of the training data.
- These subsets are created through **bootstrapping** — random sampling *with replacement*.
- The predictions from all models are then **averaged** (for regression) or **voted** (for classification).

Objective:

- **Reduce variance** (i.e., make the model more stable and less sensitive to data noise).
- Works well for **high-variance, low-bias** models like Decision Trees.

Example algorithms:

- Random Forest (ensemble of Decision Trees using bagging)
-

2. Boosting

Approach:

- Boosting builds models **sequentially**, where each new model **focuses on correcting the errors** made by the previous models.
- It assigns **higher weights** to misclassified or poorly predicted samples so that future models pay more attention to them.
- The final model is a **weighted sum** of all weak learners.

Objective:

- **Reduce bias** and improve predictive accuracy by turning weak learners into a strong learner.
- Works well for **high-bias, low-variance** problems.

Example algorithms:

- AdaBoost, Gradient Boosting, XGBoost, LightGBM, CatBoost



Summary Table:

| Feature | Bagging | Boosting |
|--------------------|------------------------------------|---------------------------------------|
| Model Training | Parallel (independent models) | Sequential (each depends on previous) |
| Data Sampling | Bootstrap samples (random subsets) | Weighted sampling based on errors |
| Objective | Reduce variance | Reduce bias |
| Model Combination | Simple averaging or voting | Weighted sum of weak learners |
| Typical Base Model | Decision Trees | Shallow Decision Trees (stumps) |
| Example Algorithms | Random Forest | AdaBoost, Gradient Boosting, XGBoost |

 **In short:**

- **Bagging = Parallel learners** → variance reduction.
- **Boosting = Sequential learners** → bias reduction.

QUESTION 2 : Explain how the Random Forest Classifier reduces overfitting compared to a single decision tree. Mention the role of two key hyperparameters in this process.

ANSWER 2 : A **Random Forest Classifier** reduces overfitting compared to a single Decision Tree by combining the predictions of **many independent trees**, each trained on a **random subset** of the data and features.

Here's how it works and why it helps:

◆ **Fundamental Idea:**

- A **single Decision Tree** tends to overfit — it learns noise and specific patterns from the training data, leading to poor generalization.
 - A **Random Forest** builds **multiple trees (an ensemble)**, each slightly different, and then **averages** (for regression) or **takes a majority vote** (for classification). This averaging process **reduces variance** and prevents overfitting.
-

◆ **How Randomness Helps:**

1. **Bootstrap Sampling (Bagging):**

- Each tree is trained on a random sample (with replacement) of the training data.
- This means each tree sees slightly different data → less correlation between trees → reduced overfitting.

2. **Random Feature Selection:**

- At each split, the model considers only a random subset of features, not all.

- This prevents strong predictors from dominating every tree and adds diversity to the ensemble.
-

◆ Two Key Hyperparameters That Help Control Overfitting:

1. **n_estimators** (Number of Trees):

- More trees → better averaging and lower variance.
- Beyond a certain point, adding more trees gives diminishing returns but still stabilizes the model.

2. **max_features** (Number of Features Considered at Each Split):

- Smaller **max_features** → trees become more diverse → less overfitting.
 - Larger **max_features** → trees become more similar → higher risk of overfitting (like a single large tree).
-

✓ Summary:

| Aspect | Decision Tree | Random Forest |
|---------------------|-------------------|--|
| Structure | Single, deep tree | Many shallow, random trees |
| Variance | High | Reduced via averaging |
| Overfitting | Common | Significantly reduced |
| Key Hyperparameters | – | n_estimators , max_features |

In short:

Random Forest reduces overfitting by using **bagging** and **feature randomness** to create diverse trees and then **aggregating** their results.

The hyperparameters **n_estimators** (number of trees) and **max_features** (feature randomness) are key to controlling this process.

QUESTION 3 : What is Stacking in ensemble learning? How does it differ from traditional bagging/boosting methods? Provide a simple example use case.

ANSWER 3 : **Stacking (Stacked Generalization)** is an **ensemble learning technique** where **multiple different models (base learners)** are trained, and then their predictions are **combined using another model (meta-learner or blender)** to make the final prediction.

◆ **How Stacking Works:**

1. **Train Base Models (Level-0 models):**

Multiple algorithms (e.g., Decision Tree, SVM, Logistic Regression) are trained on the same dataset.

2. **Generate Meta Features:**

Each base model makes predictions (on validation data). These predictions are treated as **new features**.

3. **Train Meta-Model (Level-1 model):**

A new model (often a simple one like Linear Regression or Logistic Regression) learns how to best combine the base model outputs.

◆ **Difference from Bagging and Boosting:**

| Aspect | Bagging | Boosting | Stacking |
|---------------------------|---|---|--|
| Base Learners | Same algorithm (e.g., many Decision Trees in Random Forest) | Same algorithm trained sequentially (e.g., Decision Trees in AdaBoost, XGBoost) | Different algorithms can be combined |
| Training Style | Independent and parallel | Sequential (each model fixes errors of the previous) | Parallel at base level, then combined via meta-model |
| Combination Method | Averaging (for regression) or Voting (for classification) | Weighted combination | Learned combination using a meta-model |
| Goal | Reduce variance | Reduce bias | Leverage strengths of diverse models |

◆ **Example Use Case:**

Predicting loan default (classification problem)

● **Base Models (Level-0):**

- Random Forest → captures non-linear relationships
- Logistic Regression → good for linear trends
- Gradient Boosting → strong on complex interactions

● **Meta-Model (Level-1):**

- Logistic Regression → learns how to combine predictions from the three base models optimally

Result:

The stacked model often outperforms any individual model because it integrates their diverse perspectives.

 **Summary:**

Stacking = combining multiple different models through a meta-learner to improve predictive performance.

Unlike **bagging** and **boosting**, stacking **learns how to combine** model predictions instead of just averaging or weighting them.

QUESTION 4 : What is the OOB Score in Random Forest, and why is it useful? How does it help in model evaluation without a separate validation set?

ANSWER 4 :  **OOB (Out-of-Bag) Score in Random Forest**

Definition:

The **OOB (Out-of-Bag) Score** is an internal validation method used in **Random Forests** to estimate the model's performance **without needing a separate validation or test set.**

 **How It Works**

When building each decision tree in a Random Forest:

- The algorithm uses **bootstrap sampling** — it randomly selects samples **with replacement** from the training data to train the tree.
- On average, about **63%** of the data points are selected (some are repeated).
- The remaining **~37%** of samples are **not used** to train that tree — these are called **Out-of-Bag (OOB) samples**.

After the forest is trained:

- Each data point is predicted **only by the trees that did not see it** during training (i.e., trees for which it was OOB).
 - The final **OOB prediction** for each sample is obtained by **aggregating** these predictions (majority vote for classification, average for regression).
 - The **OOB Score** is the accuracy (or R^2 for regression) computed from these OOB predictions.
-



Why It's Useful

- **✓ No need for a separate validation set:**
The OOB Score acts like an internal cross-validation, saving data for training.
 - **✓ Efficient and unbiased estimate:**
It provides an **unbiased estimate of model performance**, as each prediction comes from trees that haven't seen that data point.
 - **✓ Saves computation:**
Since it reuses training data, no extra validation process is required.
-



Example

```
from sklearn.datasets import load_iris  
from sklearn.ensemble import RandomForestClassifier  
  
# Load dataset
```

```

X, y = load_iris(return_X_y=True)

# Train Random Forest with OOB score enabled
rf = RandomForestClassifier(n_estimators=100, oob_score=True, random_state=42)
rf.fit(X, y)

# Print OOB score
print("OOB Score:", rf.oob_score_)

```

Output Example:

OOB Score: 0.96

In Summary

| Aspect | Description |
|--|---|
| Concept | Performance estimate using unseen (OOB) samples |
| Approx. size of OOB data per tree | ~37% of total data |
| Main benefit | No need for a separate validation set |
| Used for | Estimating generalization accuracy efficiently |
| Similar to | Cross-validation, but faster and built-in |

In short:

The **OOB Score** provides a **built-in, unbiased validation measure** in Random Forests, letting you assess model performance **without holding out extra data**.

QUESTION 5 : Compare AdaBoost and Gradient Boosting in terms of: • How they handle errors from weak learners • Weight adjustment mechanism • Typical use cases

ANSWER 5 : Here's a clear **comparison between AdaBoost and Gradient Boosting** across the requested points:

| Aspect | AdaBoost (Adaptive Boosting) | Gradient Boosting |
|--|--|--|
| 1. Handling errors from weak learners | Focuses more on misclassified samples from previous learners — increases their weights so the next learner pays more attention to them. | Focuses on residual errors (prediction errors) — each new learner tries to correct the residuals (differences between actual and predicted values) from the previous model. |
| 2. Weight adjustment mechanism | Adjusts sample weights after each iteration: correctly classified samples get lower weights, and misclassified ones get higher weights. | Does not modify sample weights directly; instead, it fits new learners to minimize the loss function's gradient (i.e., learns from residuals). |
| 3. Typical use cases | Works well with simple base learners (e.g., decision stumps) and when data is not too noisy. Often used for classification tasks . | More flexible and powerful — used in both regression and classification , and serves as the foundation for libraries like XGBoost, LightGBM, and CatBoost . |

Summary:

- **AdaBoost** adjusts **weights of samples** to focus on hard-to-classify points.
- **Gradient Boosting** adjusts **model predictions** iteratively to minimize overall loss using gradient descent.
- **AdaBoost** is simpler and faster; **Gradient Boosting** is more general and powerful for complex, large-scale problems.

QUESTION 6 : Why does CatBoost perform well on categorical features without requiring extensive preprocessing? Briefly explain its handling of categorical variables.

ANSWER 6 : CatBoost performs well on **categorical features** because it has a **built-in, efficient way to handle them** — without the need for manual preprocessing like one-hot encoding or label encoding.

- ◆ **Key Idea:**

CatBoost uses a technique called “**Ordered Target Statistics**” (or **Ordered Encoding**) to convert categorical values into numerical form **based on target values**, while **avoiding target leakage**.

◆ How CatBoost Handles Categorical Variables:

1. Target-Based Encoding:

Each categorical feature is transformed into a numeric value based on statistics of the target variable (e.g., mean target value for that category).

$$[\text{Encoded Value} = \frac{\text{Sum of target values for the category}}{\text{prior}} \cdot \text{Count of occurrences of the category} + 1]$$

2. Ordered Encoding to Prevent Leakage:

Instead of using the whole dataset (which could leak future information), CatBoost encodes each sample using only **data from previous samples** in a random permutation order.

3. Efficient Combination of Categorical Features:

CatBoost can also create **combinations of categorical features** automatically, capturing complex interactions that would normally require manual feature engineering.

◆ Why It Performs Well:

- No need for one-hot encoding (saves memory and time).
 - Avoids overfitting and target leakage through ordered statistics.
 - Automatically captures relationships between categorical and numerical features.
-

In short:

CatBoost natively handles categorical features by applying **ordered target-based encoding**, which makes it both **efficient and accurate** without heavy preprocessing.

QUESTION 7 : KNN Classifier Assignment: Wine Dataset Analysis with Optimization Task: 1. Load the Wine dataset (`sklearn.datasets.load_wine()`). 2. Split data into 70% train and 30% test. 3. Train a KNN classifier (default K=5) without scaling and evaluate using: a. Accuracy b. Precision, Recall, F1-Score (print classification report) 4. Apply StandardScaler, retrain KNN, and compare metrics. 5. Use GridSearchCV to find the best K (test K=1 to 20) and distance metric (Euclidean, Manhattan).

ANSWER 7 : Here's the full Python code for your **KNN Classifier Assignment on the Wine Dataset**, including all requested steps, evaluation metrics, scaling, and optimization using GridSearchCV 

```
# KNN Classifier Assignment: Wine Dataset Analysis with Optimization

from sklearn.datasets import load_wine
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, classification_report

# 1. Load the Wine dataset
wine = load_wine()
X, y = wine.data, wine.target

# 2. Split data into 70% train and 30% test
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42, stratify=y
)

# 3. Train KNN (K=5) without scaling
knn_default = KNeighborsClassifier(n_neighbors=5)
knn_default.fit(X_train, y_train)
y_pred_default = knn_default.predict(X_test)

print("== Without Scaling ==")
print("Accuracy:", accuracy_score(y_test, y_pred_default))
print("\nClassification Report:\n", classification_report(y_test, y_pred_default))

# 4. Apply StandardScaler, retrain KNN, and compare
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

knn_scaled = KNeighborsClassifier(n_neighbors=5)
```

```

knn_scaled.fit(X_train_scaled, y_train)
y_pred_scaled = knn_scaled.predict(X_test_scaled)

print("\n==== With StandardScaler ===")
print("Accuracy:", accuracy_score(y_test, y_pred_scaled))
print("\nClassification Report:\n", classification_report(y_test, y_pred_scaled))

# 5. Hyperparameter tuning using GridSearchCV
param_grid = {
    'n_neighbors': list(range(1, 21)), # K = 1 to 20
    'metric': ['euclidean', 'manhattan']
}

grid = GridSearchCV(
    KNeighborsClassifier(),
    param_grid,
    cv=5,
    scoring='accuracy',
    n_jobs=-1
)
grid.fit(X_train_scaled, y_train)

print("\n==== GridSearchCV Results ===")
print("Best Parameters:", grid.best_params_)
print("Best Cross-Validation Accuracy:", grid.best_score_)

# Evaluate best model on test data
best_knn = grid.best_estimator_
y_pred_best = best_knn.predict(X_test_scaled)

print("\n==== Optimized Model Evaluation ===")
print("Test Accuracy:", accuracy_score(y_test, y_pred_best))
print("\nClassification Report:\n", classification_report(y_test, y_pred_best))

```

Explanation of Steps

1. **Dataset** – Loads the Wine dataset from `sklearn.datasets`.
2. **Train-Test Split** – 70% for training and 30% for testing (stratified for class balance).
3. **Initial KNN (K=5)** – Trained on unscaled data; prints accuracy and classification report.

4. **Standardization** – Uses `StandardScaler` to normalize features, improving distance-based performance.
 5. **Optimization** – Uses `GridSearchCV` to tune `n_neighbors` (1–20) and distance metrics (`euclidean`, `manhattan`).
-

QUESTION 8 : PCA + KNN with Variance Analysis and Visualization Task: 1. Load the Breast Cancer dataset (`sklearn.datasets.load_breast_cancer()`). 2. Apply PCA and plot the scree plot (explained variance ratio). 3. Retain 95% variance and transform the dataset. 4. Train KNN on the original data and PCA-transformed data, then compare accuracy. 5. Visualize the first two principal components using a scatter plot (color by class).

ANSWER 8 : Here's a **complete Python program** that performs PCA + KNN with variance analysis and visualization using the **Breast Cancer dataset**:

```
# PCA + KNN with Variance Analysis and Visualization
```

```
# -----
```

```
from sklearn.datasets import load_breast_cancer
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt
import numpy as np
```

```
# 1. Load the dataset
```

```
data = load_breast_cancer()
```

```
X = data.data

y = data.target


# Standardize the data before PCA

scaler = StandardScaler()

X_scaled = scaler.fit_transform(X)


# 2. Apply PCA and plot the scree plot

pca = PCA()

pca.fit(X_scaled)

explained_variance = np.cumsum(pca.explained_variance_ratio_)


plt.figure(figsize=(8,5))

plt.plot(np.arange(1, len(explained_variance)+1), explained_variance, marker='o', color='b')

plt.title('Scree Plot (Cumulative Explained Variance)')

plt.xlabel('Number of Principal Components')

plt.ylabel('Cumulative Explained Variance Ratio')

plt.grid(True)

plt.show()


# 3. Retain 95% variance and transform the dataset

pca_95 = PCA(0.95)

X_pca = pca_95.fit_transform(X_scaled)

print(f"Number of components retained for 95% variance: {pca_95.n_components_}")
```

```
# 4. Train KNN on original and PCA-transformed data and compare accuracy

X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.3, random_state=42)

# KNN on original data

knn_original = KNeighborsClassifier(n_neighbors=5)

knn_original.fit(X_train, y_train)

y_pred_original = knn_original.predict(X_test)

acc_original = accuracy_score(y_test, y_pred_original)

# KNN on PCA data

X_train_pca, X_test_pca, _, _ = train_test_split(X_pca, y, test_size=0.3, random_state=42)

knn_pca = KNeighborsClassifier(n_neighbors=5)

knn_pca.fit(X_train_pca, y_train)

y_pred_pca = knn_pca.predict(X_test_pca)

acc_pca = accuracy_score(y_test, y_pred_pca)

print(f"Accuracy on Original Data: {acc_original:.4f}")

print(f"Accuracy on PCA (95% variance) Data: {acc_pca:.4f}")

# 5. Visualize the first two principal components

pca_2 = PCA(n_components=2)

X_2D = pca_2.fit_transform(X_scaled)
```

```
plt.figure(figsize=(8,6))

plt.scatter(X_2D[:, 0], X_2D[:, 1], c=y, cmap='coolwarm', alpha=0.7)

plt.title('PCA - First Two Principal Components')

plt.xlabel('Principal Component 1')

plt.ylabel('Principal Component 2')

plt.colorbar(label='Class (0 = Malignant, 1 = Benign)')

plt.show()
```

Explanation:

1. **Dataset:** Loads the Breast Cancer dataset from `sklearn.datasets`.
 2. **Scaling:** Standardization ensures PCA works properly (since it's variance-based).
 3. **Scree Plot:** Shows how much variance each principal component explains (helps choose optimal number).
 4. **PCA (95%):** Retains components explaining **95%** of the total variance.
 5. **KNN Comparison:**
 - Trains KNN ($k=5$) on **original scaled data** and **PCA-reduced data**.
 - Compares their **accuracy scores**.
 6. **Visualization:** Uses a **scatter plot of first two principal components** colored by class to show class separability.
-

QUESTION 9 : KNN Regressor with Distance Metrics and K-Value Analysis Task: 1. Generate a synthetic regression dataset (`sklearn.datasets.make_regression(n_samples=500, n_features=10)`). 2. Train a KNN regressor with: a. Euclidean distance ($K=5$) b. Manhattan distance ($K=5$) c. Compare Mean Squared Error (MSE) for both. 3. Test $K=1, 5, 10, 20, 50$ and plot K vs. MSE to analyze bias-variance tradeoff.

ANSWER 9 : Here's the complete **Python program** for the KNN Regressor with distance metrics and K-value analysis 

```
# KNN Regressor with Distance Metrics and K-Value Analysis
```

```
# -----
```

```
from sklearn.datasets import make_regression  
from sklearn.model_selection import train_test_split  
from sklearn.neighbors import KNeighborsRegressor  
from sklearn.metrics import mean_squared_error  
import matplotlib.pyplot as plt
```

```
# 1. Generate synthetic regression dataset
```

```
X, y = make_regression(n_samples=500, n_features=10, noise=10, random_state=42)
```

```
# Split into training and testing sets
```

```
X_train, X_test, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
# 2a. Train KNN Regressor with Euclidean distance (K=5)
```

```
knn_euclidean = KNeighborsRegressor(n_neighbors=5, metric='euclidean')
```

```
knn_euclidean.fit(X_train, y_train)
```

```
y_pred_euclidean = knn_euclidean.predict(X_test)
```

```
mse_euclidean = mean_squared_error(y_test, y_pred_euclidean)
```

```
# 2b. Train KNN Regressor with Manhattan distance (K=5)
```

```
knn_manhattan = KNeighborsRegressor(n_neighbors=5, metric='manhattan')
```

```

knn_manhattan.fit(X_train, y_train)

y_pred_manhattan = knn_manhattan.predict(X_test)

mse_manhattan = mean_squared_error(y_test, y_pred_manhattan)

print("==== Distance Metric Comparison (K=5) ====")

print(f"Euclidean Distance MSE : {mse_euclidean:.3f}")
print(f"Manhattan Distance MSE : {mse_manhattan:.3f}")

# 3. Analyze bias-variance tradeoff by varying K

k_values = [1, 5, 10, 20, 50]

mse_values = []

for k in k_values:

    knn = KNeighborsRegressor(n_neighbors=k, metric='euclidean')

    knn.fit(X_train, y_train)

    y_pred = knn.predict(X_test)

    mse = mean_squared_error(y_test, y_pred)

    mse_values.append(mse)

# Plot K vs. MSE

plt.figure(figsize=(8, 5))

plt.plot(k_values, mse_values, marker='o', linestyle='-', color='blue')

plt.title("K vs Mean Squared Error (Euclidean Distance)")

plt.xlabel("Number of Neighbors (K)")

```

```

plt.ylabel("Mean Squared Error")

plt.grid(True)

plt.show()

# Interpretation

print("\nInterpretation:")

print("- Lower K (e.g., K=1) → Low bias, high variance (fits noise).")

print("- Higher K (e.g., K=50) → High bias, low variance (smoother predictions).")

print("- Optimal K usually balances bias and variance.")

```

Explanation:

1. Dataset Generation

- Uses `make_regression()` to create a synthetic regression dataset (500 samples, 10 features).

2. Distance Metrics Comparison

- Trains two `KNeighborsRegressor` models:
 - **Euclidean distance:** default metric for continuous features.
 - **Manhattan distance:** uses absolute differences.
- Compares **Mean Squared Error (MSE)** for both.

3. Bias–Variance Analysis

- Tests `K = [1, 5, 10, 20, 50]` and plots **K vs. MSE**.
- Demonstrates how increasing `K` increases bias but reduces variance.

